

COMS20012: Integer Overflow

Joseph Hallett

bristol.ac.uk



What is all this about then?

- CPUs have registers
- Patterns in bits in registers represent stuff

```
#include <stdio.h>

int main(void) {
    unsigned int x = -1;
    void *ref = (void *)&x;

    printf("Unsigned int: %u\n", *((unsigned int *)ref));
    printf("Signed int:   %d\n", *((signed int *)ref));
    printf("Float:         %f\n", *((float *)ref));
    printf("Double:        %lf\n", *((double *)ref));
    printf("Pointer:       %p\n", *((void **)ref));

    return 0;
}
```



It's just a representation...

```
#include <stdio.h>

int main(void) {
    unsigned int x = -1;
    void *ref = (void *)&x;

    printf("Unsigned int: %u\n", *((unsigned int *)ref));
    printf("Signed int:   %d\n", *((signed int *)ref));
    printf("Float:         %f\n", *((float *)ref));
    printf("Double:        %lf\n", *((double *)ref));
    printf("Pointer:       %p\n", *((void **)ref));

    return 0;
}
```

```
Unsigned int: 4294967295
Signed int:   -1
Float:        nan
Double:       0.000000
Pointer:      0xffffffff
```



Not Quite Integers

- What a CPU (typically) regards as an integer isn't really an integer in a purely mathematical sense
- A *native* CPU integer typically has *limits* dependent on what the ISA supports
 - (some ISAs support some really weird data types... e.g. BCD)
- Operating systems again impose limits on integers
 - see `limits.h` in your OS libraries and POSIX standards



Binary and 2s Complement Notation

- *Usually* unsigned integers are represented as standard binary
 - unsigned integers typically go from 0 to $2^{\text{wordsize}}-1$
 - So for an 8 bit CPU: 0 to $2^8-1 = 255$
 - For a 32 bit CPU: 0 to $2^{32}-1 = 4,294,967,295$
- *Usually* signed integers are represented in 2s complement
 - Highest bit's value is negated
 - So $-2^{(\text{wordsize}-1)}$ to $2^{(\text{wordsize}-1)}-1$
 - So for an 8 bit CPU: -2^7 to $2^7-1 = -128$ to 127
 - For a 32 bit CPU: -2^{31} to $2^{31}-1 = -2,147,483,648$ to 2,147,483,647
- So what happens when you go beyond these limits?



What's next?

```
#include <stdio.h>

int main(void) {
    unsigned int x = -1;
    void *ref = (void *)&x;

    printf("Unsigned int: %u\n", *((unsigned int *)ref)+1);
    printf("Signed int:   %d\n", *((signed int *)ref)+1);
    printf("Float:        %f\n", *((float *)ref)+1);
    printf("Double:       %lf\n", *((double *)ref)+1);
    printf("Pointer:      %p\n", *((void **)ref)+1);

    return 0;
}
```



Wraparound and weirdness!

```
#include <stdio.h>

int main(void) {
    unsigned int x = -1;
    void *ref = (void *)&x;

    printf("Unsigned int: %u\n", *((unsigned int *)ref)+1);
    printf("Signed int:   %d\n", *((signed int *)ref)+1);
    printf("Float:        %f\n", *((float *)ref)+1);
    printf("Double:       %lf\n", *((double *)ref)+1);
    printf("Pointer:      %p\n", *((void **)ref)+1);

    return 0;
}
```

```
Unsigned int: 0
Signed int:   0
Float:        nan
Double:       1.000000
Pointer:      0x100000000
```



Whole bunch of weird rules

- Expansion from a smaller data type to a larger one *should* work
 - e.g. char to an int
- Contraction from a larger data type to a smaller one may not!
 - e.g. int to char
- Switching between signed and unsigned types...
 - Do you zero extend or one extend?



Aren't there meant to be standards?

- The C standard says what's *supposed to happen*
 - But it is full of edge cases and imprecise
- Compilers can and do differ
- ...Which means compilers support bugs from old compilers/standards to avoid breaking 50 year old programs
- It's a mess!




POLICY —

Sorry ma'am, you didn't win \$43M—there was a slot machine "malfunction"

Casino had apologized "for any inconvenience this may have caused."

DAVID KRAVETS - 6/15/2017, 2:01 PM



[Enlarge](#) / Katrina Bookman takes a selfie showing she hit the big one.

Imagine, if you would, how absolutely giddy you'd be if you won a \$43 million jackpot while playing a casino slot machine. You could burn a lot of bridges with that amount of cash.



Preventing overflows

```
#include <limits.h>
int main(void) {
    unsigned int ui;
    signed int si;
    unsigned long ul;
    signed long sl;

    ui = si = ul = sl = ULONG_MAX;

    ui = si;
    ui = ul;
    ui = sl;
    si = ui + sl;
    si = ui;
    si = ul;
    si = sl;
    ul *= ul*sl;

    return 0;
}
```

```
[$ gcc -Wall -Wextra example.c --std=c2x -pedantic
[$ # ...wat.
```

bristol.ac.uk



The C compiler doesn't have enough warnings by default

```
[$ clang -Wall -Wextra example.c --std=c2x -pedantic -Weverything
warning: include location '/usr/local/include' is unsafe for cross-compilation [-Wpoison-system-directories]
example.c:8:21: warning: implicit conversion changes signedness: 'long' to 'unsigned long' [-Wsign-conversion]
    ui = si = ul = sl = ULONG_MAX;
                        ^
note:   ^~~~~~
example.c:8:23: warning: implicit conversion changes signedness: 'unsigned long' to 'long' [-Wsign-conversion]
    ui = si = ul = sl = ULONG_MAX;
                        ^
note:   ^~~~~~
/Applications/Xcode.app/Contents/Developer/Toolchains/XcodeDefault.xctoolchain/usr/lib/clang/13.0.0/include/limits.h:57:37: note: expanded from macro 'ULONG_MAX'
#define ULONG_MAX  (__LONG_MAX__ *2UL+1UL)
                        ^
example.c:10:8: warning: implicit conversion changes signedness: 'int' to 'unsigned int' [-Wsign-conversion]
    ui = si;
        ^
example.c:14:8: warning: implicit conversion changes signedness: 'unsigned int' to 'int' [-Wsign-conversion]
    si = ui;
        ^
example.c:17:12: warning: implicit conversion changes signedness: 'long' to 'unsigned long' [-Wsign-conversion]
    ul *= ul*sl;
        ^
example.c:11:8: warning: implicit conversion loses integer precision: 'unsigned long' to 'unsigned int' [-Wshorten-64-to-32]
    ui = ul;
        ^
example.c:12:8: warning: implicit conversion loses integer precision: 'long' to 'unsigned int' [-Wshorten-64-to-32]
    ui = sl;
        ^
example.c:13:11: warning: implicit conversion loses integer precision: 'long' to 'int' [-Wshorten-64-to-32]
    si = ui + sl;
        ^
example.c:15:8: warning: implicit conversion loses integer precision: 'unsigned long' to 'int' [-Wshorten-64-to-32]
    si = ul;
        ^
example.c:16:8: warning: implicit conversion loses integer precision: 'long' to 'int' [-Wshorten-64-to-32]
    si = sl;
        ^
11 warnings generated.
```



Do your own checking

- Use the limits.h header to find limits

```
▪ unsigned int a, b;  
  if (a < UINT_MAX - b) { return a + b; }  
  else {  
    /* Look maybe you want to use a higher level  
    programming language if this is likely to be an  
    issue. There are a bunch of libraries and non-  
    standard ways of dealing with this but at the end  
    of the day you have better things to be doing  
    with your lives.  
    */  
    return -1; // ;-)
```

